

baseline

September 16, 2019

```
[1]: import gc
gc.collect()
```

[1]: 62

```
[2]: # Load libraries

import numpy as np

import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report,
    ↳ multilabel_confusion_matrix
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier

import xgboost as xgb
import lightgbm as lgb
```

```
[3]: # Load datasets

TRAIN_FILEPATH = 'data/train.csv'
train_df = pd.read_csv( TRAIN_FILEPATH, header=0 )
display(train_df.shape)
```

(15120, 56)

```
[4]: # Split training set on train/validation sets (60-20-20, 70-30)

# save target value
train_label = train_df['Cover_Type'] # {1;2;...;6;7}, 2160 entires each
train_df = train_df.drop( ['Id', 'Cover_Type'], axis=1 )

# train_test_split
```

```

VALIDATION_SIZE = 0.3

X_tr, X_val, y_tr, y_val = train_test_split(
    train_df, train_label,
    test_size=VALIDATION_SIZE,
    shuffle=True
)

```

[5]: # 1. *sklearn.LogisticRegression*, *l2 regularization*

```

logreg_l2_model = LogisticRegression(
    C=1.0,
    solver='lbfgs', # multinomial loss
    penalty='l2',
    max_iter=1000,
    multi_class='multinomial',
    n_jobs=-1
)

logreg_l2_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('logreg_model', logreg_l2_model)
])

%time logreg_l2_pipeline.fit( X_tr, y_tr )

```

CPU times: user 54.3 ms, sys: 63.5 ms, total: 118 ms
 Wall time: 8.36 s

```

[5]: Pipeline(memory=None,
      steps=[('scaler',
              StandardScaler(copy=True, with_mean=True, with_std=True)),
             ('logreg_model',
              LogisticRegression(C=1.0, class_weight=None, dual=False,
                                fit_intercept=True, intercept_scaling=1,
                                l1_ratio=None, max_iter=1000,
                                multi_class='multinomial', n_jobs=-1,
                                penalty='l2', random_state=None,
                                solver='lbfgs', tol=0.0001, verbose=0,
                                warm_start=False))],
      verbose=False)

```

[6]: %time logreg_y_val_pred = logreg_l2_pipeline.predict(X_val)

CPU times: user 16.1 ms, sys: 19.1 ms, total: 35.2 ms
 Wall time: 23.7 ms

```
[7]: display( accuracy_score(y_val, logreg_y_val_pred) )

print( classification_report(y_val, logreg_y_val_pred) )

display( multilabel_confusion_matrix(y_val, logreg_y_val_pred) )
```

0.7074514991181657

	precision	recall	f1-score	support
1	0.66	0.67	0.66	671
2	0.61	0.53	0.57	667
3	0.63	0.54	0.58	617
4	0.81	0.88	0.84	627
5	0.73	0.80	0.76	681
6	0.61	0.67	0.64	631
7	0.87	0.87	0.87	642
accuracy			0.71	4536
macro avg	0.70	0.71	0.70	4536
weighted avg	0.70	0.71	0.70	4536

```
array([[3630, 235],
       [ 222, 449]],

       [[3648, 221],
       [ 315, 352]],

       [[3726, 193],
       [ 283, 334]],

       [[3779, 130],
       [  76, 551]],

       [[3657, 198],
       [ 138, 543]],

       [[3638, 267],
       [ 210, 421]],

       [[3811,  83],
       [  83, 559]]])
```

```
[8]: # 2. sklearn.LogisticRegression, l1 regularization
```

```

logreg_l1_model = LogisticRegression(
    C=1.0,
    solver='liblinear', # one-vs-rest scheme
    penalty='l1',
    max_iter=1000,
    multi_class='ovr' # can't use 'multinomial'
)

logreg_l1_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('logreg_model', logreg_l1_model)
])

%time logreg_l1_pipeline.fit( X_tr, y_tr )

```

CPU times: user 20.6 s, sys: 54 ms, total: 20.6 s
 Wall time: 20.6 s

```

[8]: Pipeline(memory=None,
      steps=[('scaler',
              StandardScaler(copy=True, with_mean=True, with_std=True)),
             ('logreg_model',
              LogisticRegression(C=1.0, class_weight=None, dual=False,
                                fit_intercept=True, intercept_scaling=1,
                                l1_ratio=None, max_iter=1000,
                                multi_class='ovr', n_jobs=None,
                                penalty='l1', random_state=None,
                                solver='liblinear', tol=0.0001, verbose=0,
                                warm_start=False))],
      verbose=False)

```

```

[9]: %time logreg_y_val_pred = logreg_l1_pipeline.predict( X_val )

```

CPU times: user 33.6 ms, sys: 19.6 ms, total: 53.2 ms
 Wall time: 11 ms

```

[10]: display( accuracy_score(y_val, logreg_y_val_pred) )

print( classification_report(y_val, logreg_y_val_pred) )

display( multilabel_confusion_matrix(y_val, logreg_y_val_pred) )

```

0.6783509700176367

precision	recall	f1-score	support
-----------	--------	----------	---------

1	0.64	0.62	0.63	671
2	0.58	0.49	0.53	667
3	0.60	0.54	0.57	617
4	0.80	0.89	0.84	627
5	0.64	0.72	0.68	681
6	0.60	0.62	0.61	631
7	0.87	0.87	0.87	642
accuracy				0.68 4536
macro avg	0.67	0.68	0.68	4536
weighted avg	0.67	0.68	0.67	4536

```
array([[3629, 236],
       [ 252, 419]],

       [[3634, 235],
       [ 341, 326]],

       [[3697, 222],
       [ 284, 333]],

       [[3770, 139],
       [  68, 559]],

       [[3576, 279],
       [ 191, 490]],

       [[3643, 262],
       [ 242, 389]],

       [[3808,  86],
       [  81, 561]]])
```

[27]: *# 3. sklearn.RandomForestClassifier*

```
rfc_model = RandomForestClassifier(
    n_estimators=1000,
    n_jobs=-1
)

%time rfc_model.fit( X_tr, y_tr )
```

CPU times: user 23.4 s, sys: 380 ms, total: 23.8 s
Wall time: 3.66 s

```
[27]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                             max_depth=None, max_features='auto', max_leaf_nodes=None,
                             min_impurity_decrease=0.0, min_impurity_split=None,
                             min_samples_leaf=1, min_samples_split=2,
                             min_weight_fraction_leaf=0.0, n_estimators=1000,
                             n_jobs=-1, oob_score=False, random_state=None, verbose=0,
                             warm_start=False)
```

```
[28]: %time rfc_y_val_pred = rfc_model.predict( X_val )
```

CPU times: user 1.56 s, sys: 60.2 ms, total: 1.62 s
Wall time: 407 ms

```
[29]: display( accuracy_score(y_val, rfc_y_val_pred) )

print( classification_report(y_val, rfc_y_val_pred) )

display( multilabel_confusion_matrix(y_val, rfc_y_val_pred) )
```

0.8619929453262787

	precision	recall	f1-score	support
1	0.81	0.75	0.78	671
2	0.79	0.72	0.75	667
3	0.87	0.79	0.83	617
4	0.92	0.97	0.95	627
5	0.88	0.95	0.91	681
6	0.83	0.90	0.86	631
7	0.93	0.96	0.95	642
accuracy			0.86	4536
macro avg	0.86	0.86	0.86	4536
weighted avg	0.86	0.86	0.86	4536

```
array([[3744, 121],
       [ 168, 503]],

      [[3739, 130],
       [ 187, 480]],

      [[3844, 75],
       [ 129, 488]],

      [[3856, 53],
       [ 16, 611]],
```

```

[[3767, 88],
 [ 37, 644]],

[[3790, 115],
 [ 65, 566]],

[[3850, 44],
 [ 24, 618]]])

```

```

[14]: display(
      pd.DataFrame({
          'feature_name': X_tr.columns,
          'feature_imp': rfc_model.feature_importances_
      }).sort_values( by='feature_imp', ascending=False ).head()
    )

```

	feature_name	feature_imp
0	Elevation	0.231019
5	Horizontal_Distance_To_Roadways	0.088981
9	Horizontal_Distance_To_Fire_Points	0.070936
3	Horizontal_Distance_To_Hydrology	0.063224
4	Vertical_Distance_To_Hydrology	0.054716

```

[15]: # 4. xgb.XGBClassifier

eval_set = [ (X_val, y_val) ]

xgb_model = xgb.XGBClassifier(
    gamma=0.025,
    learning_rate=0.35,
    max_depth=5,
    n_estimators=1000,
    objective='multi:softmax',
    n_jobs=4
)

%time xgb_model.fit( X_tr, y_tr, eval_set=eval_set, eval_metric='merror',
    verbose=False )

```

CPU times: user 3min 1s, sys: 495 ms, total: 3min 1s
 Wall time: 45.7 s

```

[15]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                  colsample_bynode=1, colsample_bytree=1, gamma=0.025,

```

```

learning_rate=0.35, max_delta_step=0, max_depth=5,
min_child_weight=1, missing=None, n_estimators=1000, n_jobs=4,
nthread=None, objective='multi:softprob', random_state=0,
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,
silent=None, subsample=1, verbosity=1)

```

```
[16]: %time xgb_y_val_pred = xgb_model.predict( X_val )
```

```

CPU times: user 1.33 s, sys: 4.03 ms, total: 1.33 s
Wall time: 335 ms

```

```
[17]: display( accuracy_score(y_val, xgb_y_val_pred) )

print( classification_report(y_val, xgb_y_val_pred) )

display( multilabel_confusion_matrix(y_val, xgb_y_val_pred) )
```

```
0.8538359788359788
```

	precision	recall	f1-score	support
1	0.77	0.71	0.74	671
2	0.73	0.69	0.71	667
3	0.85	0.83	0.84	617
4	0.95	0.96	0.96	627
5	0.89	0.94	0.92	681
6	0.84	0.89	0.86	631
7	0.93	0.96	0.95	642
accuracy			0.85	4536
macro avg	0.85	0.86	0.85	4536
weighted avg	0.85	0.85	0.85	4536

```

array([[3721, 144],
       [ 194, 477]],

       [[3700, 169],
        [ 208, 459]],

       [[3831,  88],
        [ 107, 510]],

       [[3880,  29],
        [  22, 605]],

       [[3776,  79],

```



```

[ 38, 643]],

[[3795, 110],
 [ 70, 561]],

[[3850, 44],
 [ 24, 618]]])

```

```

[18]: display(
      pd.DataFrame({
          'feature_name': X_tr.columns,
          'feature_imp': xgb_model.feature_importances_
      }).sort_values( by='feature_imp', ascending=False ).head()
    )

```

	feature_name	feature_imp
13	Wilderness_Area4	0.081424
43	Soil_Type30	0.076421
25	Soil_Type12	0.065155
16	Soil_Type3	0.062327
0	Elevation	0.059893

```

[19]: # 5. lgb.LGBMClassifier

lgb_model = lgb.LGBMClassifier(
    learning_rate=0.2,
    max_depth=-1,
    n_estimators=1000,
    objective='multiclass',
    n_jobs=8,
    verbose=0
)

%time lgb_model.fit( X_tr, y_tr )

```

CPU times: user 1min 27s, sys: 607 ms, total: 1min 27s
 Wall time: 11.7 s

```

[19]: LGBMClassifier(boosting_type='gbdt', class_weight=None, colsample_bytree=1.0,
    importance_type='split', learning_rate=0.2, max_depth=-1,
    min_child_samples=20, min_child_weight=0.001, min_split_gain=0.0,
    n_estimators=1000, n_jobs=8, num_leaves=31,
    objective='multiclass', random_state=None, reg_alpha=0.0,
    reg_lambda=0.0, silent=True, subsample=1.0,
    subsample_for_bin=200000, subsample_freq=0, verbose=0)

```

```
[20]: %time lgb_y_val_pred = lgb_model.predict( X_val )
```

CPU times: user 4.75 s, sys: 9 μ s, total: 4.75 s
Wall time: 625 ms

```
[21]: display( accuracy_score(y_val, lgb_y_val_pred) )  
  
print( classification_report(y_val, lgb_y_val_pred) )  
  
display( multilabel_confusion_matrix(y_val, lgb_y_val_pred) )
```

0.8699294532627866

	precision	recall	f1-score	support
1	0.79	0.74	0.77	671
2	0.77	0.71	0.74	667
3	0.87	0.84	0.86	617
4	0.96	0.98	0.97	627
5	0.89	0.96	0.92	681
6	0.85	0.91	0.88	631
7	0.94	0.97	0.95	642
accuracy			0.87	4536
macro avg	0.87	0.87	0.87	4536
weighted avg	0.87	0.87	0.87	4536

```
array([[3735, 130],  
       [ 173, 498]],  
  
       [[3725, 144],  
       [ 191, 476]],  
  
       [[3845, 74],  
       [ 101, 516]],  
  
       [[3883, 26],  
       [ 15, 612]],  
  
       [[3777, 78],  
       [ 29, 652]],  
  
       [[3804, 101],  
       [ 59, 572]],
```

```
[[3857, 37],
 [ 22, 620]])
```

```
[22]: display(
      pd.DataFrame({
          'feature_name': X_tr.columns,
          'feature_imp': lgb_model.feature_importances_
      }).sort_values( by='feature_imp', ascending=False ).head()
    )
```

	feature_name	feature_imp
5	Horizontal_Distance_To_Roadways	23481
9	Horizontal_Distance_To_Fire_Points	23437
0	Elevation	22266
4	Vertical_Distance_To_Hydrology	15688
3	Horizontal_Distance_To_Hydrology	14445

```
[23]: #####
```

```
[33]: # Blend baseline LGBclf, XGBclf and sklearn RandomForestClassifier models into
      → submission
```

```
# Load test set
TEST_FILEPATH = 'data/test.csv'
test_df = pd.read_csv( TEST_FILEPATH, header=0 )
display(test_df.shape)

# Save 'Id' for submission
test_ids = test_df['Id']
test_df = test_df.drop( ['Id'], axis=1 )
```

```
(565892, 55)
```

```
[41]: print('random forest classifier...')
      rfc_model = RandomForestClassifier(
          n_estimators=1000,
          n_jobs=-1
      )
      rfc_model.fit( train_df, train_label )
      # rfc_pred = rfc_model.predict( test_df )

      print('xgboost classifier...')
      xgb_model = xgb.XGBClassifier(
          gamma=0.025,
          learning_rate=0.35,
```

```

        max_depth=9,
        n_estimators=1000,
        objective='multi:softmax',
        n_jobs=4
    )
    xgb_model.fit( train_df, train_label )
    # xgb_pred = xgb_model.predict( test_df )

    print('lgbm classifier...')
    lgb_model = lgb.LGBMClassifier(
        learning_rate=0.2,
        max_depth=-1,
        n_estimators=1000,
        objective='multiclass',
        n_jobs=8,
        verbose=0
    )
    lgb_model.fit( train_df, train_label )
    # lgb_pred = lgb_model.predict( test_df )

```

```

[45]: rfc_pred_proba = rfc_model.predict_proba( test_df )
      xgb_pred_proba = xgb_model.predict_proba( test_df )
      lgb_pred_proba = lgb_model.predict_proba( test_df )

```

```

[46]: blended_pred_proba = ( rfc_pred_proba + xgb_pred_proba + lgb_pred_proba ) / 3.0

```

```

[78]: # display(blended_pred_proba.shape, blended_pred_proba)

```

```

def get_argmax( x_array ):
    return np.argmax( x_array )
get_argmax_vect = np.vectorize(get_argmax)

y_pred_max_proba = np.array( [get_argmax(x_row)+1 for x_row in
    ↪blended_pred_proba] )

```

```

[84]: submission = pd.DataFrame(
      y_pred_max_proba, index=test_ids, columns=['Cover_Type']
    )
    submission.to_csv('submission_baseline_rfc_lgb_xgb.csv', index_label='Id')

```

```

[ ]: # Kaggle public leaderboard score (score - classification accuracy): 0.77121

```