

feateng_model

September 17, 2019

```
[17]: # Import libraries

import re

%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns

import numpy as np

import pandas as pd

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import multilabel_confusion_matrix, accuracy_score, \
    →classification_report

from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestClassifier

import xgboost as xgb
import lightgbm as lgb
```

```
[2]: # Load datasets

TRAIN_FILEPATH = 'data/train.csv'
train_df = pd.read_csv( TRAIN_FILEPATH, header=0 )
display(train_df.shape)

TEST_FILEPATH = 'data/test.csv'
test_df = pd.read_csv( TEST_FILEPATH, header=0 )
display(test_df.shape)
```

(15120, 56)

(565892, 55)

```

[3]: # Fix values

# 1. As we can see, some of 'Hillshade_3pm' values = 0 -> fix this
train_df['Hillshade_9am'].hist(bins=200, label='9am')
train_df['Hillshade_3pm'].hist(bins=200, label='3pm')
train_df['Hillshade_Noon'].hist(bins=200, label='Noon')
plt.title('train before')
plt.legend()
plt.show()

display(
    'train before:',
    train_df[ train_df['Hillshade_3pm'] < 5 ]['Hillshade_3pm'].value_counts()
)

test_df['Hillshade_9am'].hist(bins=200, label='9am')
test_df['Hillshade_3pm'].hist(bins=200, label='3pm')
test_df['Hillshade_Noon'].hist(bins=200, label='Noon')
plt.title('test before')
plt.legend()
plt.show()

display(
    'test before',
    test_df[ test_df['Hillshade_3pm'] < 5 ]['Hillshade_3pm'].value_counts()
)

# Use data from test set (500k rows), forget about train set (15k)
hillshade_pred_useful_columns = [
    col_name for col_name
    in test_df.columns.values
    if col_name not in ['Hillshade_3pm', 'Id', 'Cover_Type']
]
train_topredict_rows = train_df[ train_df['Hillshade_3pm'] == 0 ]
    ↳ [hillshade_pred_useful_columns]
test_topredict_rows = test_df[ test_df['Hillshade_3pm'] == 0 ]
    ↳ [hillshade_pred_useful_columns]

totrain_rows_X = test_df[ test_df['Hillshade_3pm'] != 0 ]
    ↳ [hillshade_pred_useful_columns]
totrain_rows_y = test_df[ test_df['Hillshade_3pm'] != 0 ]['Hillshade_3pm']

predictor = KNeighborsRegressor(n_neighbors=10) # todo: estimate accuracy
predictor.fit( totrain_rows_X, totrain_rows_y )

train_df.loc[ train_topredict_rows.index.values, 'Hillshade_3pm' ] = np.around(
    ↳ predictor.predict(train_topredict_rows) )

```

```

test_df.loc[ test_topredict_rows.index.values, 'Hillshade_3pm' ] = np.around(
    →predictor.predict(test_topredict_rows) )

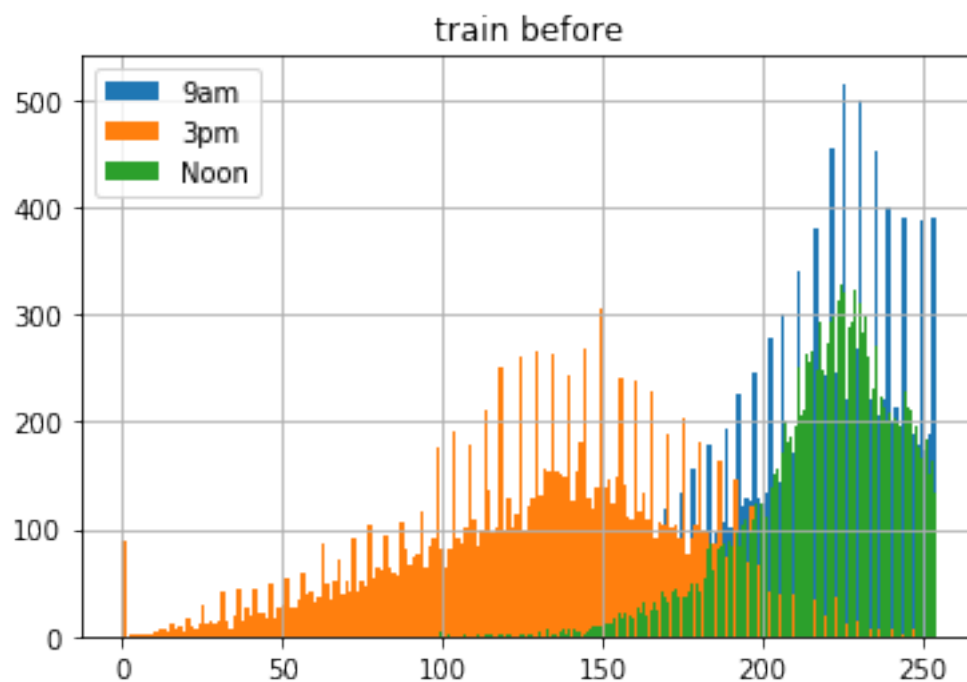
# 'after' result :
plt.title('train after')
train_df['Hillshade_9am'].hist(bins=200, label='9am')
train_df['Hillshade_3pm'].hist(bins=200, label='3pm')
train_df['Hillshade_Noon'].hist(bins=200, label='Noon')
plt.legend()
plt.show()

display(
    'train after:',
    train_df[ train_df['Hillshade_3pm'] < 5 ]['Hillshade_3pm'].value_counts()
)

plt.title('test after')
test_df['Hillshade_9am'].hist(bins=200, label='9am')
test_df['Hillshade_3pm'].hist(bins=200, label='3pm')
test_df['Hillshade_Noon'].hist(bins=200, label='Noon')
plt.legend()
plt.show()

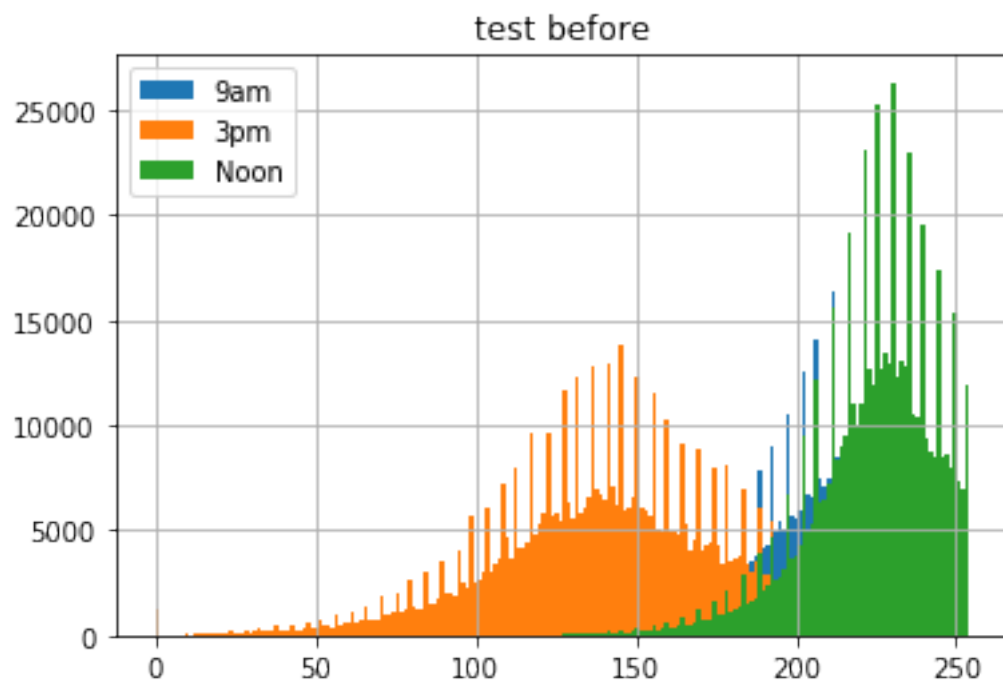
display(
    'test after',
    test_df[ test_df['Hillshade_3pm'] < 5 ]['Hillshade_3pm'].value_counts()
)

```




'train before:'

```
0    88
3     3
4     1
1     1
Name: Hillshade_3pm, dtype: int64
```



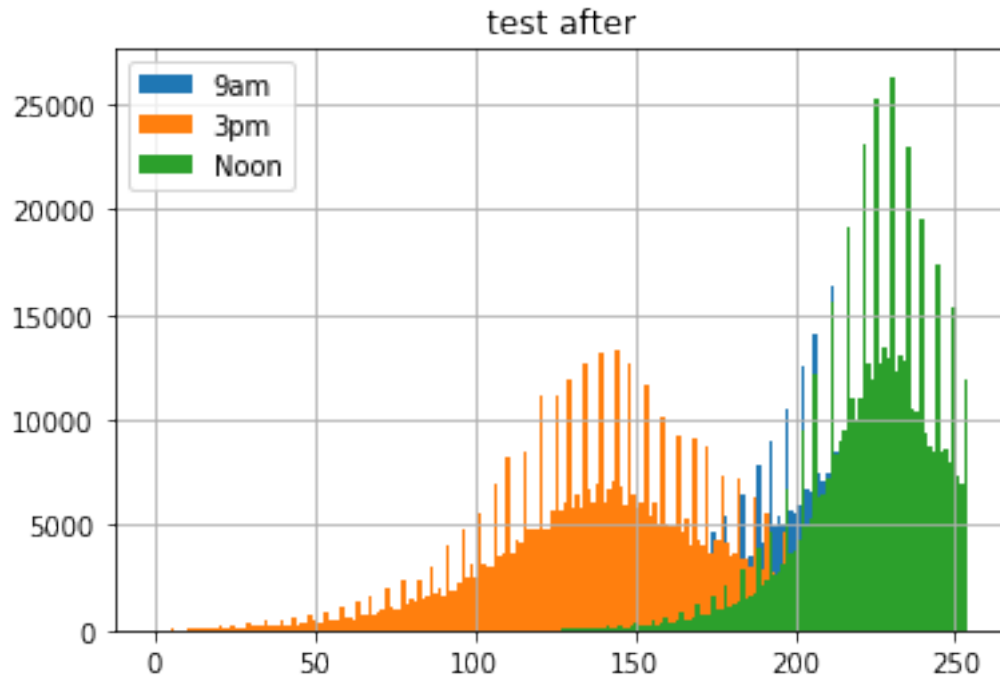
'test before'

```
0    1250
4     19
2     15
1     14
3     12
Name: Hillshade_3pm, dtype: int64
```



'train after:'

```
3.0    3
4.0    1
1.0    1
Name: Hillshade_3pm, dtype: int64
```



'test after'

```
4.0    19
2.0    15
1.0    14
3.0    12
Name: Hillshade_3pm, dtype: int64
```

[4]: *# Work with concatenated feautures*

```
traintest_df = pd.concat( [train_df, test_df], sort=False, ignore_index=True )
```

[5]: *# Combine Soil_TypeX and Wilderness_AreaX into single feature*

```
def merge_onehot( dataset_df, col_name_no_x ):
    """Convert col_nameX features to single feature
    X means some integer value.
    Doesn't work with multiple calls - returns 0s for all col_name_no_x.
    """
    dataset_df_cpy = dataset_df.copy()
    # 1. Identify columns
    all_df_columns = dataset_df_cpy.columns.values
    re_pattern_compiled = re.compile( "~{0}(\d+)$".format( col_name_no_x ) )
```

```

matched_columns = list(filter( re_pattern_compiled.match, all_df_columns ))
# 2. Change columns: multiply by 'x' value
for matched_column in matched_columns:
    col_name_x_value = re_pattern_compiled.match(matched_column).groups()[0]
    dataset_df_cpy[matched_column] *= int( col_name_x_value )
# 3. Merge col_name_x columns into single col_name column
dataset_df_cpy[col_name_no_x] = 0
for matched_column in matched_columns:
    dataset_df_cpy[col_name_no_x] += dataset_df_cpy[matched_column]
# 4. Drop col_name_x columns
dataset_df_cpy = dataset_df_cpy.drop( matched_columns, axis=1 )
return dataset_df_cpy

def _ugly_merge_wildernessarea_soiltype_traintest( train_or_test_df ):
    train_or_test_df = merge_onehot( train_or_test_df,
    →col_name_no_x='Wilderness_Area' )
    train_or_test_df = merge_onehot( train_or_test_df,
    →col_name_no_x='Soil_Type' )
    return train_or_test_df

# train_df = _ugly_merge_wildernessarea_soiltype_traintest( train_df )
# test_df = _ugly_merge_wildernessarea_soiltype_traintest( test_df )

# display(train_df.shape, test_df.shape)

traintest_df = _ugly_merge_wildernessarea_soiltype_traintest( traintest_df )

```

[6]: # Feature engineering

```

def add_soil_family_inplace( df ):
    # src: https://www.kaggle.com/c/forest-cover-type-prediction/data
    # (soiltypeX, soiltypeY, ...): family_name_str
    soiltype_family_mapping = {
        (2, 4): 1, # 'ratake',
        (10, 11, 13, 32, 33): 2, # 'catamount',
        (21, 22, 23, 24, 25, 27, 28): 3, # 'leighcan',
        (38, 39, 40): 4, # 'moran'
    #     (... all other IDs ...): 'other_type'
    }
    df['soil_family'] = 5 # 'other_type'
    for i in df.index:
        soil_type_value = df.at[i, 'Soil_Type']
        for key in soiltype_family_mapping.keys():
            if soil_type_value in key:
                df.at[i, 'soil_family'] = soiltype_family_mapping[key]

```

```

def add_soil_complex_inplace( df ):
    # (soiltypeX, soiltypeY, ...): complex_name_str
    soiltype_complex_mapping = {
        (1, 3, 4, 5, 6, 10, 27, 28, 33): 1, # 'rock_outcrop',
        (11, 12, 34, 40): 2, # 'rock_land',
        (20, 23): 3, # 'typic_cryaquolls',
        (26, 31): 4, # 'catamaount_families',
        (29, 30): 5, # 'legault_family',
        (32, 39): 6, # 'leighcan_family',
    }
    df['soil_complex'] = 7 # 'other_type'
    for i in df.index:
        soil_type_value = df.at[i, 'Soil_Type']
        for key in soiltype_complex_mapping.keys():
            if soil_type_value in key:
                df.at[i, 'soil_complex'] = soiltype_complex_mapping[key]

def add_soil_stonetype_inplace( df ):
    # (soiltypeX, soiltypeY, ...): stonetype_name_str
    soiltype_stonetype_mapping = {
        (1, 2, 6, 9, 12, 18, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 36, 37, 38, 39, 40): 1, # 'stony',
        (3, 4, 5, 10, 11, 13, ): 2, # 'rubbly'
    }
    df['soil_stonetype'] = 3 # 'other_type'
    for i in df.index:
        soil_type_value = df.at[i, 'Soil_Type']
        for key in soiltype_stonetype_mapping.keys():
            if soil_type_value in key:
                df.at[i, 'soil_stonetype'] = soiltype_stonetype_mapping[key]

def main_feature_engineering( df ):
    df_cpy = df.copy()

    # Median hillshade index [0-255]
    df_cpy['median_hillshade_idx'] = df_cpy[['Hillshade_9am', 'Hillshade_Noon', 'Hillshade_3pm']].agg('median', axis='columns')

    # Mean hillshade index [0-255]
    df_cpy['median_hillshade_idx'] = df_cpy[['Hillshade_9am', 'Hillshade_Noon', 'Hillshade_3pm']].agg('mean', axis='columns')

    # Distance to hydrology (using HorizontalDistance and VerticalDistance to hydrology) - pythagoras theorem

```



```

    df_cpy['hydrology_distance'] = np.sqrt(
→df_cpy['Horizontal_Distance_To_Hydrology']**2 +
→df_cpy['Vertical_Distance_To_Hydrology']**2 )
    # Aspect binning: 20 intervals
    ASPECT_BINS_CNT = 20
    df_cpy['aspect_bin'] = pd.cut( df_cpy['Aspect'], bins=ASPECT_BINS_CNT )
    df_cpy['aspect_bin'] = df_cpy['aspect_bin'].apply(
        lambda interval: '{0}_{1}'.format(interval.left, interval.right)
    )
    # Polynomial features: dependance of 9am->noon->3pm->9am
    df_cpy['9am_noon_dep'] = df_cpy['Hillshade_9am'] * df_cpy['Hillshade_Noon']
    df_cpy['noon_3pm_dep'] = df_cpy['Hillshade_Noon'] * df_cpy['Hillshade_3pm']
    df_cpy['3pm_9am_dep'] = df_cpy['Hillshade_3pm'] * df_cpy['Hillshade_9am']
    # Cosine of slope: relationships between hillshade and other features
    df_cpy['slope_cosine'] = np.cos( df_cpy['Slope'] )
    # Log-transform 'Elevation' feature
    df_cpy['Elevation'] = np.log1p( df_cpy['Elevation'] )

    return df_cpy

def apply_feature_engineering( df ):
    print('soil family...')
    add_soil_family_inplace( df )
    print('soil complex...')
    add_soil_complex_inplace( df )
    print('soil stonetype...')
    add_soil_stonetype_inplace( df )
    print('main feature engineering...')
    df = main_feature_engineering( df )
    return df

# train_df = apply_feature_engineering( train_df )
# test_df = apply_feature_engineering( test_df )

traintest_df = apply_feature_engineering( traintest_df )

```

```

soil family...
soil complex...
soil stonetype...
main feature engineering...

```

```

[7]: # 29th - x3
      # 15 - only 3 values

```

```
# display(
#     traintest_df['Soil_Type'].value_counts()
# )
# traintest_df.shape
```

[8]: # Convert 'aspect_bin' to numerical format

```
aspectbin_lblencoder = LabelEncoder()
traintest_df['aspect_bin'] = aspectbin_lblencoder.fit_transform(
    traintest_df['aspect_bin'] )
```

[9]: # Split traintest df for model validation

```
train = traintest_df.iloc[:train_df.shape[0], :]
test = traintest_df.iloc[train_df.shape[0]:, :]

# remove redundant columns
train_labels = train['Cover_Type']
train = train.drop( ['Id', 'Cover_Type'], axis=1 )
test = test.drop( ['Id'], axis=1 )

# train-validation split
VALIDATION_SIZE = 0.3
X_tr, X_val, y_tr, y_val = train_test_split(
    train, train_labels,
    test_size=VALIDATION_SIZE,
    shuffle=True
)

display(X_tr.shape, X_val.shape)
```

(10584, 22)

(4536, 22)

[10]: # 1. Try out lgbm

```
lgb_model = lgb.LGBMClassifier(
    learning_rate=0.25,
    max_depth=-1,
    n_estimators=1000,
    objective='multiclass',
    n_jobs=8,
    verbose=1
)
lgb_model.fit( X_tr, y_tr )
lgb_y_val_pred = lgb_model.predict( X_val )
```

```
display( accuracy_score(y_val, lgb_y_val_pred) )
print( classification_report(y_val, lgb_y_val_pred) )
display( multilabel_confusion_matrix(y_val, lgb_y_val_pred) )
```

0.8608906525573192

	precision	recall	f1-score	support
1.0	0.77	0.71	0.74	632
2.0	0.74	0.70	0.72	638
3.0	0.85	0.83	0.84	613
4.0	0.96	0.96	0.96	667
5.0	0.91	0.94	0.93	686
6.0	0.84	0.90	0.86	645
7.0	0.93	0.98	0.95	655
accuracy			0.86	4536
macro avg	0.86	0.86	0.86	4536
weighted avg	0.86	0.86	0.86	4536

```
array([[3772, 132],
       [ 182, 450]],

      [[3741, 157],
       [ 192, 446]],

      [[3835, 88],
       [ 104, 509]],

      [[3839, 30],
       [ 29, 638]],

      [[3789, 61],
       [ 42, 644]],

      [[3777, 114],
       [ 67, 578]],

      [[3832, 49],
       [ 15, 640]]])
```

```
[11]: display(
        pd.DataFrame({
            'feature_name': X_tr.columns.values,
```

```

        'feature_imp': lgb_model.feature_importances_
    }).sort_values(by='feature_imp', ascending=False)
)

```

	feature_name	feature_imp
9	Horizontal_Distance_To_Fire_Points	16296
5	Horizontal_Distance_To_Roadways	16148
0	Elevation	14868
4	Vertical_Distance_To_Hydrology	9797
16	hydrology_distance	7317
18	9am_noon_dep	7295
1	Aspect	6448
3	Horizontal_Distance_To_Hydrology	5268
6	Hillshade_9am	5238
20	3pm_9am_dep	4951
21	slope_cosine	4615
19	noon_3pm_dep	4501
11	Soil_Type	4478
15	median_hillshade_idx	4457
8	Hillshade_3pm	4255
7	Hillshade_Noon	4034
2	Slope	3358
17	aspect_bin	1974
13	soil_complex	1968
10	Wilderness_Area	1759
12	soil_family	1565
14	soil_stonetype	891

[12]: # 2. Try out xgb

```

xgb_model = xgb.XGBClassifier(
    gamma=0.03,
    learning_rate=0.2,
    max_depth=5,
    n_estimators=1000,
    objective='multi:softmax',
    n_jobs=4
)
xgb_model.fit( X_tr, y_tr )
xgb_y_val_pred = xgb_model.predict( X_val )

display( accuracy_score(y_val, xgb_y_val_pred) )
print( classification_report(y_val, xgb_y_val_pred) )
display( multilabel_confusion_matrix(y_val, xgb_y_val_pred) )

```

0.8474426807760141

	precision	recall	f1-score	support
1.0	0.76	0.70	0.73	632
2.0	0.72	0.68	0.70	638
3.0	0.83	0.81	0.82	613
4.0	0.95	0.96	0.95	667
5.0	0.91	0.93	0.92	686
6.0	0.81	0.87	0.84	645
7.0	0.94	0.97	0.95	655
accuracy			0.85	4536
macro avg	0.84	0.84	0.84	4536
weighted avg	0.84	0.85	0.85	4536

```
array([[3766, 138],
       [ 191, 441]],

      [[3724, 174],
       [ 201, 437]],

      [[3819, 104],
       [ 119, 494]],

      [[3832, 37],
       [ 29, 638]],

      [[3784, 66],
       [ 49, 637]],

      [[3762, 129],
       [ 86, 559]],

      [[3837, 44],
       [ 17, 638]]])
```

```
[13]: display(
      pd.DataFrame({
          'feature_name': X_tr.columns.values,
          'feature_imp': xgb_model.feature_importances_
      }).sort_values(by='feature_imp', ascending=False)
    )
```

	feature_name	feature_imp
11	Soil_Type	0.206622
0	Elevation	0.180959
10	Wilderness_Area	0.108293

14	soil_stonetype	0.071709
12	soil_family	0.063648
18	9am_noon_dep	0.043666
3	Horizontal_Distance_To_Hydrology	0.037654
5	Horizontal_Distance_To_Roadways	0.027700
9	Horizontal_Distance_To_Fire_Points	0.025955
13	soil_complex	0.025126
6	Hillshade_9am	0.024648
16	hydrology_distance	0.022914
7	Hillshade_Noon	0.022648
1	Aspect	0.018521
8	Hillshade_3pm	0.018179
4	Vertical_Distance_To_Hydrology	0.018095
2	Slope	0.017150
19	noon_3pm_dep	0.014901
20	3pm_9am_dep	0.014727
17	aspect_bin	0.013576
15	median_hillshade_idx	0.013010
21	slope_cosine	0.010297

[14]: # 3. Try out rfc

```
rfc_model = RandomForestClassifier(
    n_estimators=1000,
    n_jobs=-1
)
rfc_model.fit( X_tr, y_tr )
rfc_y_val_pred = rfc_model.predict( X_val )

display( accuracy_score(y_val, rfc_y_val_pred) )
print( classification_report(y_val, rfc_y_val_pred) )
display( multilabel_confusion_matrix(y_val, rfc_y_val_pred) )
```

0.8441358024691358

	precision	recall	f1-score	support
1.0	0.75	0.72	0.73	632
2.0	0.72	0.68	0.70	638
3.0	0.83	0.79	0.81	613
4.0	0.93	0.96	0.94	667
5.0	0.91	0.92	0.91	686
6.0	0.81	0.87	0.84	645
7.0	0.92	0.96	0.94	655
accuracy			0.84	4536

macro avg	0.84	0.84	0.84	4536
weighted avg	0.84	0.84	0.84	4536

```
array([[3756, 148],
       [ 180, 452]],

      [[3733, 165],
       [ 206, 432]],

      [[3824, 99],
       [ 129, 484]],

      [[3818, 51],
       [ 26, 641]],

      [[3787, 63],
       [ 56, 630]],

      [[3763, 128],
       [ 86, 559]],

      [[3828, 53],
       [ 24, 631]]])
```

```
[15]: display(
      pd.DataFrame({
          'feature_name': X_tr.columns.values,
          'feature_imp': rfc_model.feature_importances_
      }).sort_values(by='feature_imp', ascending=False)
    )
```

	feature_name	feature_imp
0	Elevation	0.207208
11	Soil_Type	0.112166
5	Horizontal_Distance_To_Roadways	0.073421
9	Horizontal_Distance_To_Fire_Points	0.057405
10	Wilderness_Area	0.054237
16	hydrology_distance	0.044296
12	soil_family	0.042523
18	9am_noon_dep	0.040009
3	Horizontal_Distance_To_Hydrology	0.037526
4	Vertical_Distance_To_Hydrology	0.036029
6	Hillshade_9am	0.033354
13	soil_complex	0.033267
1	Aspect	0.030202
8	Hillshade_3pm	0.027031

19	noon_3pm_dep	0.026675
20	3pm_9am_dep	0.025613
7	Hillshade_Noon	0.024566
15	median_hillshade_idx	0.023852
2	Slope	0.020532
14	soil_stonetype	0.020438
21	slope_cosine	0.015050
17	aspect_bin	0.014602

[16]: *# Try to increase models performance by fixing skewness*

```
# tofix_skew_col_names = []
# for col_name in traintest_df:
#     skew_value = traintest_df[col_name].skew()
#     if not -1 < skew_value < 1:
#         tofix_skew_col_names.append( col_name )

# Cannot apply boxcox1p for all columns
# for col_name in tofix_skew_col_names:
#     try:
#         boxcox_norm = boxcox_normmax( trainteset_df[col_name] + 1 )
#         display( col_name, boxcox1p(traintest_df[col_name], boxcox_norm).
→skew() )
#     except:
#         display('cannot apply boxcox for {0}'.format(col_name))

# cant do that
```

[24]: *# Find best hyperparameters for RFC, XGB, LGB classifiers*

```
# 1. RFC
rfc_model = RandomForestClassifier( )
rfc_param_grid = {
    'n_estimators': [100, 250, 500, 1000],
    'criterion': ['gini', 'entropy'],
    'max_depth': [3, 4, 5, 10, None],
    'min_samples_split': np.linspace(0.1, 1, 5),
    'max_features': [2, 5, 'auto'],
    'n_jobs': [4],
}
rfc_grid_search = GridSearchCV(
    estimator=rfc_model,
    param_grid=rfc_param_grid,
    cv=5,
    verbose=2, iid=False, n_jobs=4
)
rfc_grid_search.fit( X_tr, y_tr )
```



```
display( rfc_grid_search.best_params_ )
```

Fitting 5 folds for each of 600 candidates, totalling 3000 fits

```
[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done 33 tasks      | elapsed: 11.9s
[Parallel(n_jobs=4)]: Done 154 tasks     | elapsed: 56.2s
[Parallel(n_jobs=4)]: Done 357 tasks     | elapsed: 2.1min
[Parallel(n_jobs=4)]: Done 640 tasks     | elapsed: 4.0min
[Parallel(n_jobs=4)]: Done 1005 tasks    | elapsed: 6.1min
[Parallel(n_jobs=4)]: Done 1450 tasks    | elapsed: 9.1min
[Parallel(n_jobs=4)]: Done 1977 tasks    | elapsed: 12.6min
[Parallel(n_jobs=4)]: Done 2584 tasks    | elapsed: 16.9min
[Parallel(n_jobs=4)]: Done 3000 out of 3000 | elapsed: 19.9min finished
```

```
{'criterion': 'gini',
 'max_depth': None,
 'max_features': 5,
 'min_samples_split': 0.1,
 'n_estimators': 500,
 'n_jobs': 4}
```

[37]: # 2. LGBM

```
lgb_model = lgb.LGBMClassifier(
    objective='multiclass',
    n_jobs=4, verbose=0
)
lgb_grid_params = {
    'learning_rate': [0.2, 0.25, 0.3],
    'num_leaves': [ int(val) for val in np.linspace(5, 25, 3) ],
    'max_depth': [-1, 5, 15, 25],
    'n_estimators': [100, 250, 500, 1000],
    'min_split_gain': [0.0, 0.05, 0.1]
}
lgb_grid_search = GridSearchCV(
    estimator=lgb_model,
    param_grid=lgb_grid_params,
    cv=2,
    verbose=2, iid=False, n_jobs=4
)
lgb_grid_search.fit( X_tr, y_tr )
display( lgb_grid_search.best_params_ )
```

Fitting 2 folds for each of 432 candidates, totalling 864 fits

```
[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=4)]: Done 33 tasks      | elapsed: 12.1min
[Parallel(n_jobs=4)]: Done 154 tasks     | elapsed: 66.0min
[Parallel(n_jobs=4)]: Done 357 tasks     | elapsed: 148.4min
[Parallel(n_jobs=4)]: Done 640 tasks     | elapsed: 259.8min
[Parallel(n_jobs=4)]: Done 864 out of 864 | elapsed: 344.6min finished
```

```
{'learning_rate': 0.3,
 'max_depth': -1,
 'min_split_gain': 0.0,
 'n_estimators': 250,
 'num_leaves': 25}
```

```
[ ]: # 3. XGB

xgb_model = xgb.XGBClassifier(
    objective='multi:softmax',
    n_jobs=4, verbosity=0
)
xgb_grid_params = {
    'gamma': [0.01, 0.05, 0.1, 0.2],
    'max_depth': [3, 5, 25],
    'n_estimators': [250, 500],
}
xgb_grid_search = GridSearchCV(
    estimator=xgb_model,
    param_grid=xgb_grid_params,
    cv=2,
    verbose=2, iid=False, n_jobs=4
)
xgb_grid_search.fit( X_tr, y_tr )
display( xgb_grid_search.best_params_ )
```

Fitting 2 folds for each of 24 candidates, totalling 48 fits

```
[Parallel(n_jobs=4)]: Using backend LokyBackend with 4 concurrent workers.
```